

project-source-code

September 3, 2020

1 Project Overview

- Traffic sign classification is an important task for self driving cars.
- In this project, we will train a Deep Convolutional Neural Network (DCNN) to classify traffic sign images.
- The dataset for experiments contains 51,839 color images classified into 43 different traffic signs.
- The following is the label number with the corresponding category name:
 - 0 = Speed limit (20km/h)
 - 1 = Speed limit (30km/h)
 - 2 = Speed limit (50km/h)
 - 3 = Speed limit (60km/h)
 - 4 = Speed limit (70km/h)
 - 5 = Speed limit (80km/h)
 - 6 = End of speed limit (80km/h)
 - 7 = Speed limit (100km/h)
 - 8 = Speed limit (120km/h)
 - 9 = No passing
 - 10 = No passing for vehicles over 3.5 metric tons
 - 11 = Right-of-way at the next intersection
 - 12 = Priority road
 - 13 = Yield
 - 14 = Stop
 - 15 = No vehicles
 - 16 = Vehicles over 3.5 metric tons prohibited
 - 17 = No entry
 - 18 = General caution
 - 19 = Dangerous curve to the left
 - 20 = Dangerous curve to the right
 - 21 = Double curve
 - 22 = Bumpy road
 - 23 = Slippery road
 - 24 = Road narrows on the right
 - 25 = Road work
 - 26 = Traffic signals
 - 27 = Pedestrians

- 28 = Children crossing
- 29 = Bicycles crossing
- 30 = Beware of ice/snow
- 31 = Wild animals crossing
- 32 = End of all speed and passing limits
- 33 = Turn right ahead
- 34 = Turn left ahead
- 35 = Ahead only
- 36 = Go straight or right
- 37 = Go straight or left
- 38 = Keep right
- 39 = Keep left
- 40 = Roundabout mandatory
- 41 = End of no passing
- 42 = End of no passing by vehicles over 3.5 metric tons

```
[1]: import numpy as np
import pandas as pd
import pickle, random
from tensorflow import keras
from matplotlib import pyplot as plt
```

2 Data Extraction

```
[2]: with open("./data/train.p", mode='rb') as tr_data:
    for_training = pickle.load(tr_data)
with open("./data/valid.p", mode='rb') as val_data:
    for_validation = pickle.load(val_data)
with open("./data/test.p", mode='rb') as test_data:
    for_testing = pickle.load(test_data)
```

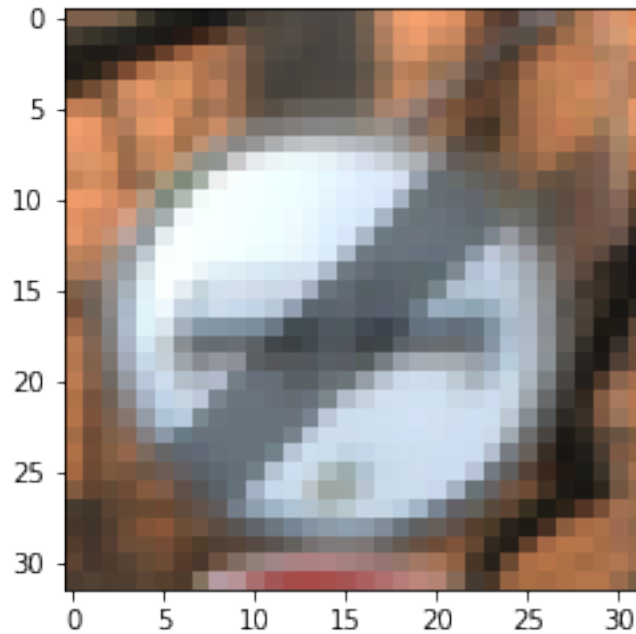
We want to separate the data into three parts that can be used training, validation, and test phases.

```
[3]: X_train, y_train = for_training['features'], for_training['labels']
X_valid, y_valid = for_validation['features'], for_validation['labels']
X_test, y_test = for_testing['features'], for_testing['labels']
```

After extracting the data, we can plot a random image and print the corresponding label number to re-check the correctness of our actions.

```
[4]: random_index = np.random.randint(0, 1000)
plt.imshow(X_train[random_index].squeeze())
print(f"The printed image is labelled as: {y_train[random_index]}")
```

The printed image is labelled as: 41



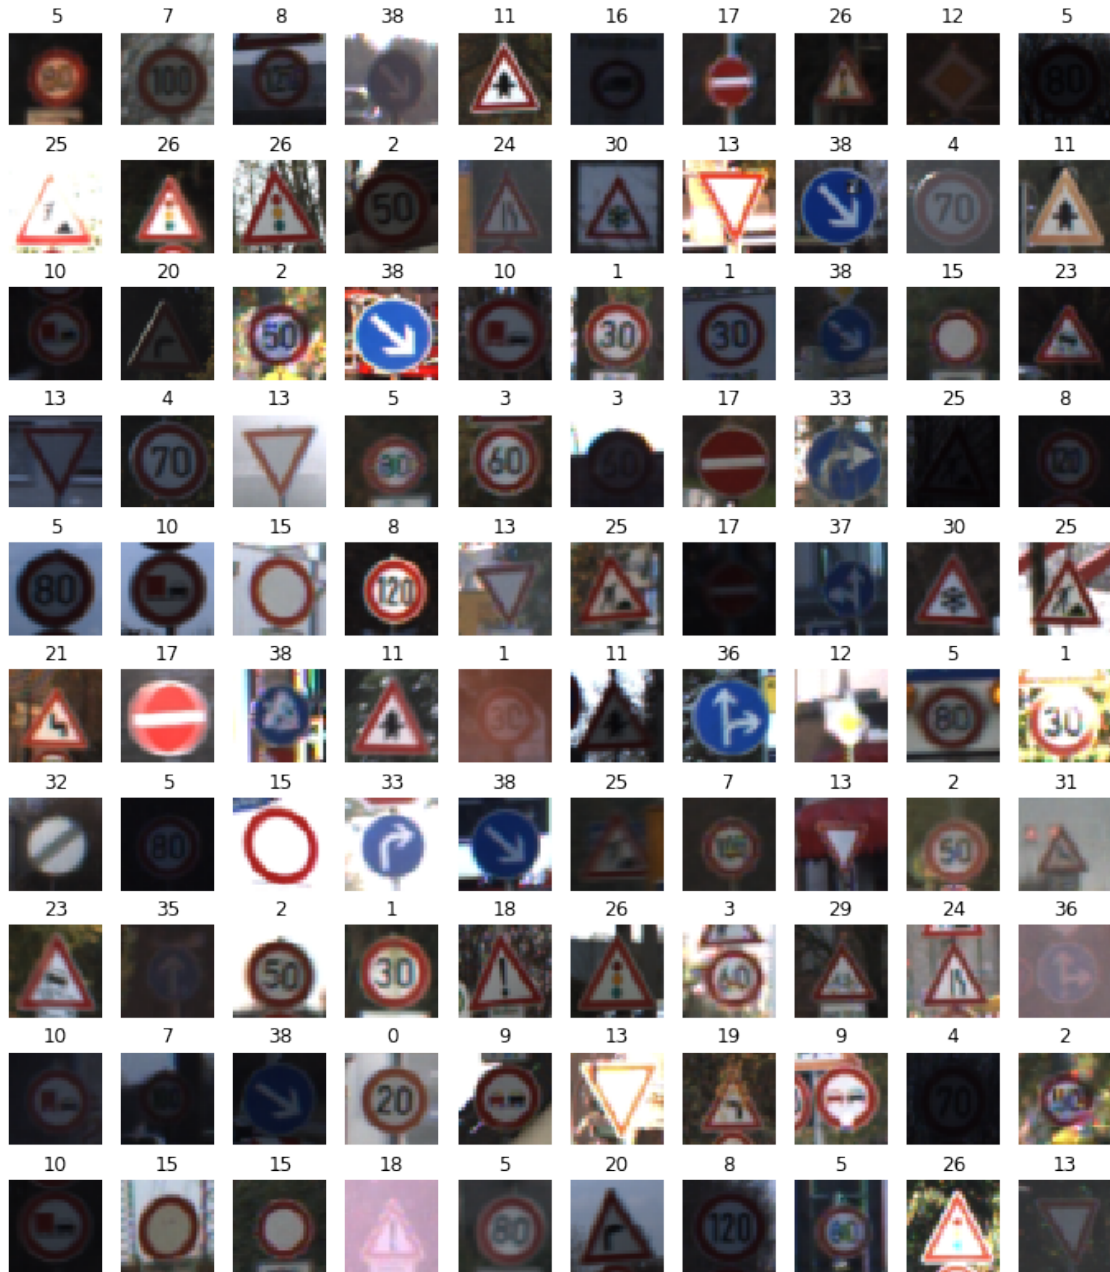
We can also plot grid of images and their corresponding labels.

```
[5]: n_rows, n_cols = 10, 10

fig, axes = plt.subplots(nrows=n_rows, ncols=n_cols, figsize=(12,14))
axes = axes.ravel() # flatten the axes: before: 10x10; after: 100
n_train = len(X_train)

for index in range(n_rows * n_cols):

    random_index = np.random.randint(0, n_train)
    axes[index].imshow(X_train[random_index].squeeze())
    axes[index].set_title(y_train[random_index])
    axes[index].axis('off')
```



3 Data pre-processing

In this step, we would like to convert the color images into grayscale. Because the color is not the decisive factor in classification of traffic signs. Also, grayscale images contain three times fewer features, which makes the training process less time consuming and less computationally expensive.

```
[6]: print(f"Dataset containing color images: {X_train.shape}")
X_train_gray = np.sum(X_train, axis=3, keepdims=True)
print(f"Dataset containing grayscale images: {X_train_gray.shape}")
```

```
Dataset containing color images: (34799, 32, 32, 3)
Dataset containing grayscale images: (34799, 32, 32, 1)
```

```
[7]: print(f"Dataset containing color images: {X_valid.shape}")
X_valid_gray = np.sum(X_valid, axis=3, keepdims=True)
print(f"Dataset containing grayscale images: {X_valid_gray.shape}")
```

```
Dataset containing color images: (4410, 32, 32, 3)
Dataset containing grayscale images: (4410, 32, 32, 1)
```

```
[8]: print(f"Dataset containing color images: {X_test.shape}")
X_test_gray = np.sum(X_test, axis=3, keepdims=True)
print(f"Dataset containing grayscale images: {X_test_gray.shape}")
```

```
Dataset containing color images: (12630, 32, 32, 3)
Dataset containing grayscale images: (12630, 32, 32, 1)
```

Next, we want to standardize the training, validation, and test data. Since, currently, the features have values ranging from 0 to 255 that makes the training and inference process slow and inefficient. Therefore, we transform the values into considerably smaller ones and make the features follow standard normal distribution. This can be achieved by subtracting mean of the training data and dividing by the standard deviation of the data from each sample in training, validation, and test data, respectively. * It is important to use mean and standard deviation of training data for the standardization. This allows to keep the data distribution the same for validation and test data, too. Otherwise, if training and validation or test data distributions are different, then the training will be useless.

```
[9]: train_mean, train_std = np.mean(X_train_gray), np.std(X_train_gray)
def standardization(data, mean, std):
    return (data - mean) / std
def stats(data):
    print(f"Mean of the given data is: {np.mean(data)}\nStd of the given data_
    ↳is: {np.std(data)}")
```

```
[10]: stats(X_train_gray)
X_train_gray_standardized = standardization(X_train_gray, train_mean, train_std)
stats(X_train_gray_standardized)
```

```
Mean of the given data is: 248.03276711098917
Std of the given data is: 198.02938725647743
Mean of the given data is: 3.162312313596526e-17
Std of the given data is: 0.9999999999999993
```

```
[11]: stats(X_valid_gray)
      X_valid_gray_standardized = standardization(X_valid_gray, train_mean, train_std)
      stats(X_valid_gray_standardized)
```

Mean of the given data is: 250.66928212691326
Std of the given data is: 203.96106434118835
Mean of the given data is: 0.013313756369449377
Std of the given data is: 1.0299535193583595

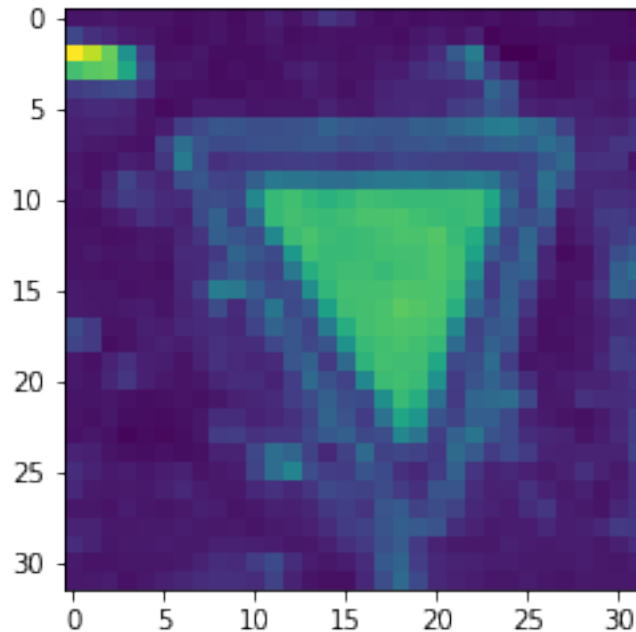
```
[12]: stats(X_test_gray)
      X_test_gray_standardized = standardization(X_test_gray, train_mean, train_std)
      stats(X_test_gray_standardized)
```

Mean of the given data is: 246.4453810836055
Std of the given data is: 200.29273072777218
Mean of the given data is: -0.00801591142292316
Std of the given data is: 1.0114293312858829

Now, we can plot a random image from any of the datasets to see difference of the above transformations.

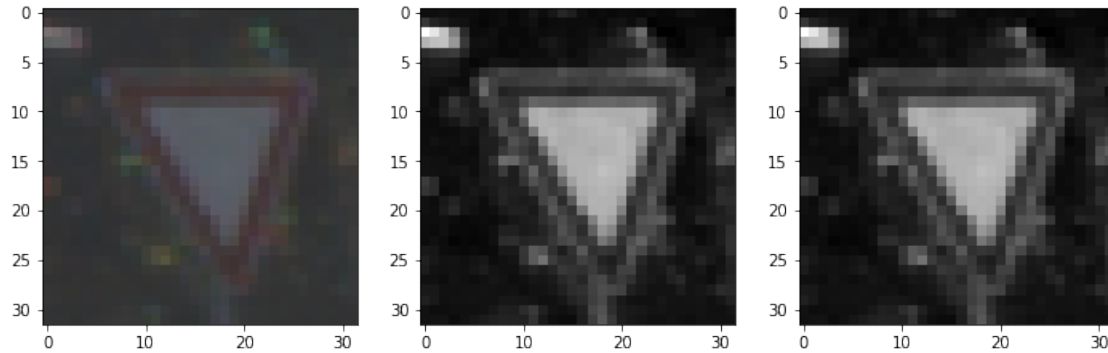
```
[13]: plt.imshow(X_train[random_index].squeeze())
      plt.imshow(X_train_gray[random_index].squeeze())
      plt.imshow(X_train_gray_standardized[random_index].squeeze())
```

```
[13]: <matplotlib.image.AxesImage at 0x217c3c2d5f8>
```



```
[14]: fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(12,5))
axes[0].imshow(X_train[random_index].squeeze())
axes[1].imshow(X_train_gray[random_index].squeeze(), cmap='gray')
axes[2].imshow(X_train_gray_standardized[random_index].squeeze(), cmap='gray')
```

```
[14]: <matplotlib.image.AxesImage at 0x217c3ccacc0>
```



4 DCNN formulation

After the data is ready for being trained, we can formulate DCNN for training the classifier.

```
[15]: model = keras.models.Sequential([
    keras.layers.Conv2D(filters=16, kernel_size=(5, 5), activation='relu',
    ↪kernel_initializer='he_normal', input_shape=(32, 32, 1)),
    keras.layers.MaxPooling2D(pool_size=2, strides=2), keras.layers.
    ↪Dropout(rate=0.2),
    keras.layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu',
    ↪kernel_initializer='he_normal', input_shape=(32, 32, 1)),
    keras.layers.MaxPooling2D(pool_size=2, strides=2), keras.layers.
    ↪Dropout(rate=0.3),
    keras.layers.Flatten(), keras.layers.Dense(units=120, activation='relu'),
    ↪keras.layers.Dense(units=43, activation='softmax')
])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #

conv2d (Conv2D)	(None, 28, 28, 16)	416

max_pooling2d (MaxPooling2D)	(None, 14, 14, 16)	0

dropout (Dropout)	(None, 14, 14, 16)	0

conv2d_1 (Conv2D)	(None, 12, 12, 32)	4640

max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 32)	0

dropout_1 (Dropout)	(None, 6, 6, 32)	0

flatten (Flatten)	(None, 1152)	0

dense (Dense)	(None, 120)	138360

dense_1 (Dense)	(None, 43)	5203
=====		
Total params: 148,619		
Trainable params: 148,619		
Non-trainable params: 0		

```
[16]: model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',  
    ↪ metrics=['accuracy'])
```

We can also use EarlyStopping Callback in order not to waste time waiting the end of the training phase. We can set the limit we want and the model training will be stopped automatically.

```
[17]: class Callback(keras.callbacks.Callback):  
  
    def on_epoch_end(self, epoch, logs={}):  
  
        if (logs.get('loss') < 0.01):  
            print("Training loss reached its limit! Cancelling training!")  
            self.model.stop_training = True  
        elif (logs.get('accuracy') > 0.995):  
            print("Training accuracy reached its limit! Cancelling training!")  
            self.model.stop_training = True  
  
callback = Callback()
```

```
[18]: results = model.fit(X_train_gray_standardized, y_train,  
    ↪ validation_data=(X_valid_gray_standardized, y_valid),  
        batch_size=256, epochs=50, verbose=1, callbacks=[callback])
```

Train on 34799 samples, validate on 4410 samples

Epoch 1/50

34799/34799 [=====] - 4s 109us/sample - loss: 2.6885 -
accuracy: 0.3150 - val_loss: 1.5617 - val_accuracy: 0.5748

Epoch 2/50

34799/34799 [=====] - 1s 35us/sample - loss: 0.9842 -
accuracy: 0.7188 - val_loss: 0.7015 - val_accuracy: 0.8190

Epoch 3/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.5216 -
accuracy: 0.8537 - val_loss: 0.4972 - val_accuracy: 0.8678

Epoch 4/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.3599 -
accuracy: 0.8990 - val_loss: 0.3883 - val_accuracy: 0.8966

Epoch 5/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.2796 -
accuracy: 0.9236 - val_loss: 0.3521 - val_accuracy: 0.9061

Epoch 6/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.2246 -
accuracy: 0.9382 - val_loss: 0.3288 - val_accuracy: 0.9098

Epoch 7/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.1886 -
accuracy: 0.9475 - val_loss: 0.3161 - val_accuracy: 0.9156

Epoch 8/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.1605 -
accuracy: 0.9560 - val_loss: 0.3168 - val_accuracy: 0.9186

Epoch 9/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.1455 -
accuracy: 0.9586 - val_loss: 0.2646 - val_accuracy: 0.9345

Epoch 10/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.1291 -
accuracy: 0.9643 - val_loss: 0.2738 - val_accuracy: 0.9299

Epoch 11/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.1144 -
accuracy: 0.9682 - val_loss: 0.2834 - val_accuracy: 0.9290

Epoch 12/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.1067 -
accuracy: 0.9706 - val_loss: 0.2390 - val_accuracy: 0.9365

Epoch 13/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0929 -
accuracy: 0.9738 - val_loss: 0.2436 - val_accuracy: 0.9363

Epoch 14/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0877 -
accuracy: 0.9755 - val_loss: 0.2389 - val_accuracy: 0.9429

Epoch 15/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0813 -
accuracy: 0.9763 - val_loss: 0.2939 - val_accuracy: 0.9329

Epoch 16/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0762 -
accuracy: 0.9785 - val_loss: 0.2741 - val_accuracy: 0.9395

Epoch 17/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0717 -
accuracy: 0.9797 - val_loss: 0.2327 - val_accuracy: 0.9438

Epoch 18/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0702 -
accuracy: 0.9790 - val_loss: 0.2377 - val_accuracy: 0.9444

Epoch 19/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0612 -
accuracy: 0.9819 - val_loss: 0.2522 - val_accuracy: 0.9431

Epoch 20/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0637 -
accuracy: 0.9820 - val_loss: 0.2004 - val_accuracy: 0.9508

Epoch 21/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0546 -
accuracy: 0.9836 - val_loss: 0.2689 - val_accuracy: 0.9451

Epoch 22/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.0574 -
accuracy: 0.9828 - val_loss: 0.2377 - val_accuracy: 0.9410

Epoch 23/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.0528 -
accuracy: 0.9846 - val_loss: 0.2240 - val_accuracy: 0.9510

Epoch 24/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0505 -
accuracy: 0.9852 - val_loss: 0.2239 - val_accuracy: 0.9465

Epoch 25/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0477 -
accuracy: 0.9862 - val_loss: 0.2199 - val_accuracy: 0.9515

Epoch 26/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0460 -
accuracy: 0.9871 - val_loss: 0.2148 - val_accuracy: 0.9503

Epoch 27/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0427 -
accuracy: 0.9871 - val_loss: 0.1971 - val_accuracy: 0.9553

Epoch 28/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0389 -
accuracy: 0.9884 - val_loss: 0.1765 - val_accuracy: 0.9585

Epoch 29/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0394 -
accuracy: 0.9874 - val_loss: 0.1897 - val_accuracy: 0.9592

Epoch 30/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0368 -
accuracy: 0.9886 - val_loss: 0.1881 - val_accuracy: 0.9535

Epoch 31/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0392 -
accuracy: 0.9881 - val_loss: 0.1887 - val_accuracy: 0.9592

Epoch 32/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0363 -
accuracy: 0.9896 - val_loss: 0.2217 - val_accuracy: 0.9551

Epoch 33/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0341 -
accuracy: 0.9890 - val_loss: 0.2163 - val_accuracy: 0.9531

Epoch 34/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.0372 -
accuracy: 0.9890 - val_loss: 0.2231 - val_accuracy: 0.9476

Epoch 35/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.0314 -
accuracy: 0.9906 - val_loss: 0.2104 - val_accuracy: 0.9574

Epoch 36/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0340 -
accuracy: 0.9896 - val_loss: 0.1951 - val_accuracy: 0.9587

Epoch 37/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.0343 -
accuracy: 0.9892 - val_loss: 0.2431 - val_accuracy: 0.9467

Epoch 38/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.0285 -
accuracy: 0.9909 - val_loss: 0.1923 - val_accuracy: 0.9580

Epoch 39/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.0275 -
accuracy: 0.9919 - val_loss: 0.2061 - val_accuracy: 0.9544

Epoch 40/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0275 -
accuracy: 0.9920 - val_loss: 0.2108 - val_accuracy: 0.9512

Epoch 41/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0276 -
accuracy: 0.9922 - val_loss: 0.2128 - val_accuracy: 0.9546

Epoch 42/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0269 -
accuracy: 0.9918 - val_loss: 0.2202 - val_accuracy: 0.9540

Epoch 43/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0264 -
accuracy: 0.9922 - val_loss: 0.2023 - val_accuracy: 0.9560

Epoch 44/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0251 -
accuracy: 0.9926 - val_loss: 0.2631 - val_accuracy: 0.9494

Epoch 45/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.0267 -
accuracy: 0.9924 - val_loss: 0.2215 - val_accuracy: 0.9533

Epoch 46/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.0204 -
accuracy: 0.9938 - val_loss: 0.2299 - val_accuracy: 0.9580

Epoch 47/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0234 -
accuracy: 0.9930 - val_loss: 0.2106 - val_accuracy: 0.9590

Epoch 48/50
34799/34799 [=====] - 1s 34us/sample - loss: 0.0257 -
accuracy: 0.9930 - val_loss: 0.1894 - val_accuracy: 0.9617

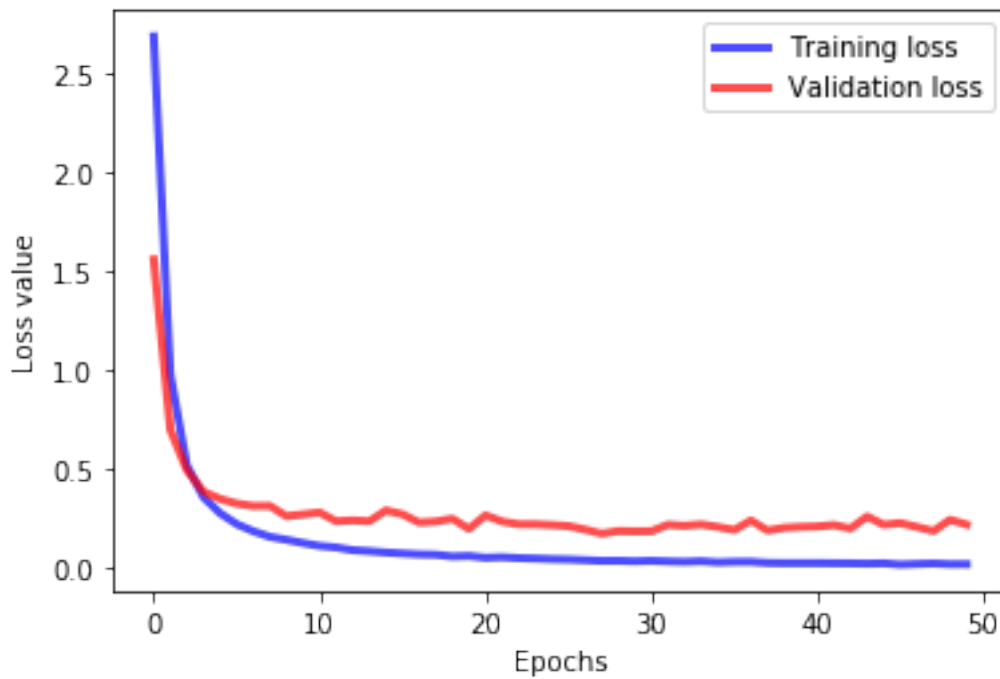
Epoch 49/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0221 -
accuracy: 0.9936 - val_loss: 0.2467 - val_accuracy: 0.9537

Epoch 50/50
34799/34799 [=====] - 1s 35us/sample - loss: 0.0223 -
accuracy: 0.9932 - val_loss: 0.2224 - val_accuracy: 0.9565

```
[19]: tr_loss = results.history['loss']
      val_loss = results.history['val_loss']
      tr_acc = results.history['accuracy']
      val_acc = results.history['val_accuracy']
      epochs = range(len(tr_loss))
```

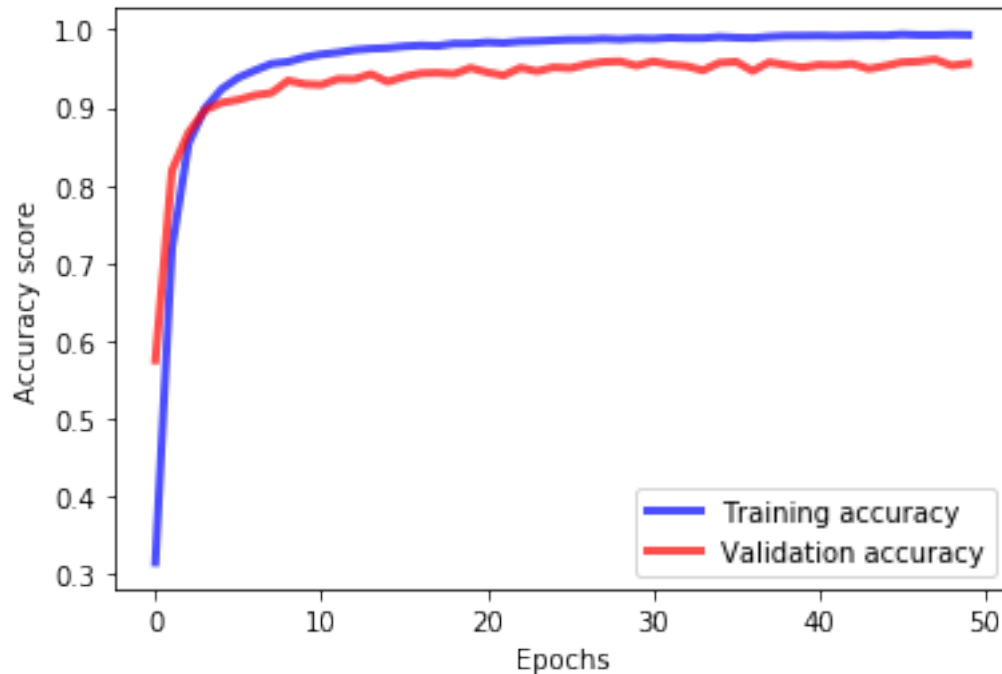
```
[20]: plt.plot(epochs, tr_loss, linewidth=3, color='b', alpha=0.7, label='Training_
      ↪loss')
      plt.plot(epochs, val_loss, linewidth=3, color='r', alpha=0.7, label='Validation_
      ↪loss')
      plt.legend()
      plt.xlabel("Epochs")
      plt.ylabel("Loss value")
```

```
[20]: Text(0, 0.5, 'Loss value')
```



```
[21]: plt.plot(epochs, tr_acc, linewidth=3, color='b', alpha=0.7, label='Training_
      ↪accuracy')
      plt.plot(epochs, val_acc, linewidth=3, color='r', alpha=0.7, label='Validation_
      ↪accuracy')
      plt.legend()
      plt.xlabel("Epochs")
      plt.ylabel("Accuracy score")
```

```
[21]: Text(0, 0.5, 'Accuracy score')
```



5 Model Evaluation

We can check the generalizability of the model by using unseen test data. Since we have the labels of the test data, we may be interested in the accuracy score of the model in test data, too. However, majority of the test datasets do not have label, so this process is not always possible.

```
[22]: from sklearn.metrics import confusion_matrix, accuracy_score

test_preds = np.argmax(model.predict(X_test_gray_standardized), axis=1)
acc_score = accuracy_score(y_test, test_preds)
print(f"Accuracy score of the model on test data: {acc_score:.3f}")
```

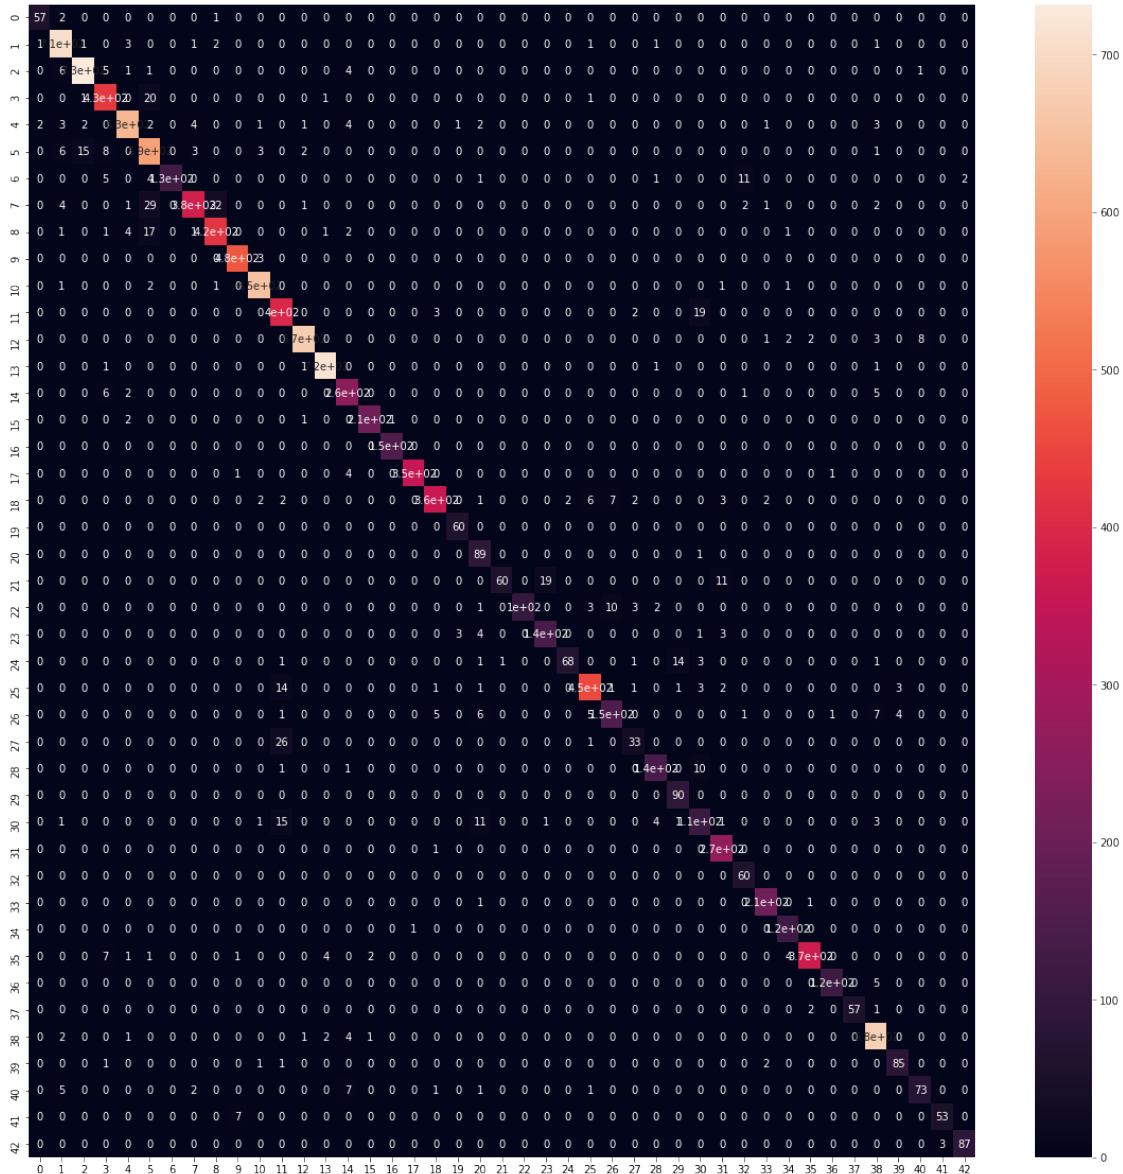
Accuracy score of the model on test data: 0.949

We may also want to plot confusion matrix to see in which particular categories the model made mistakes.

```
[23]: import seaborn as sns

cm = confusion_matrix(y_test, test_preds)
plt.figure(figsize=(20, 20))
sns.heatmap(cm, annot=True)
```

```
[23]: <matplotlib.axes._subplots.AxesSubplot at 0x218b6538748>
```



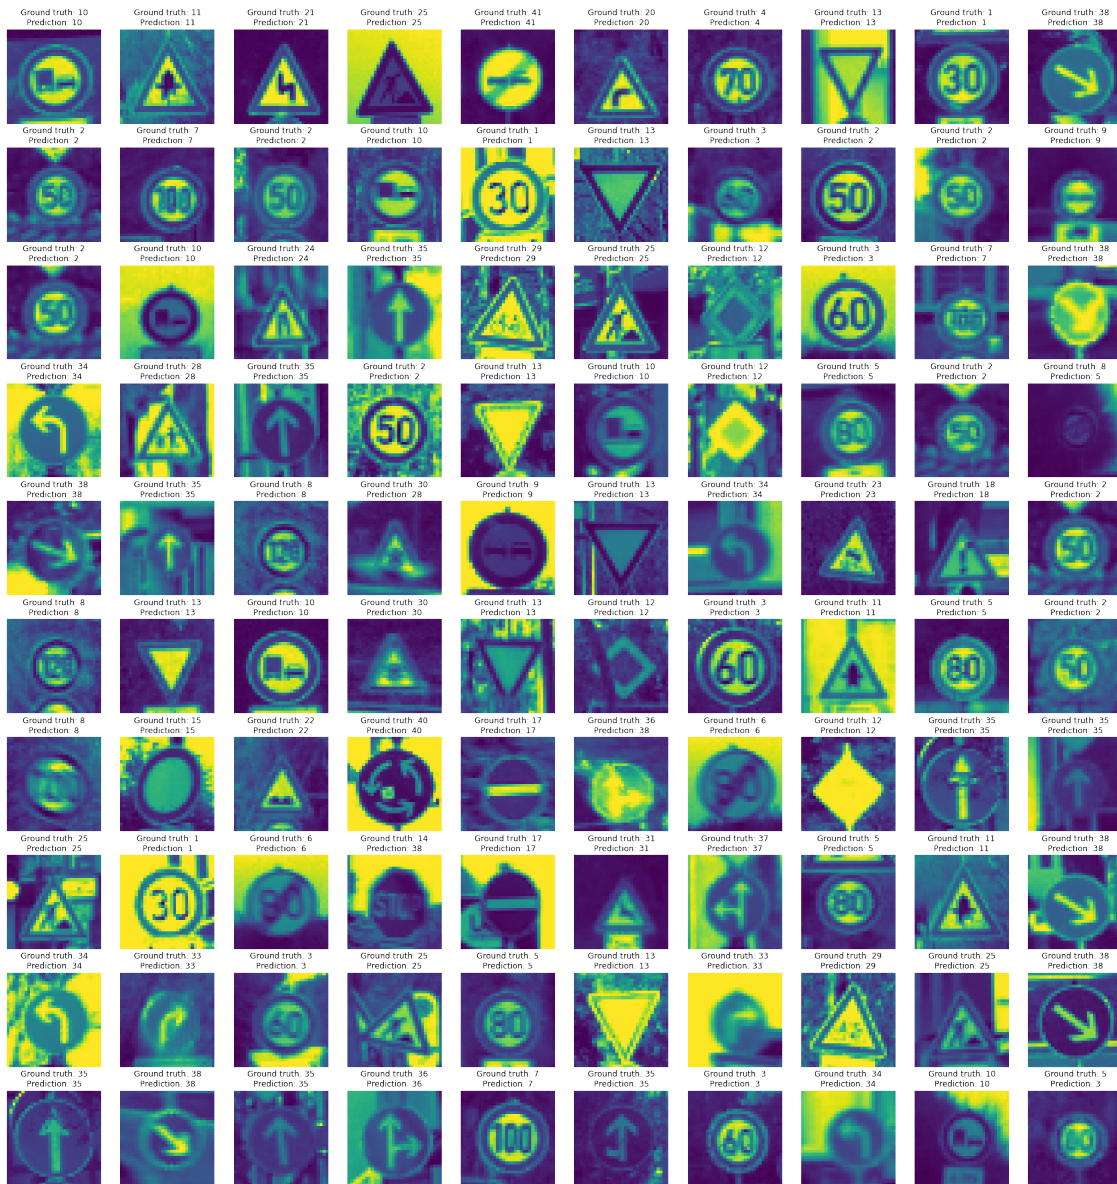
Finally, we can plot some random images from test data with their corresponding labels and observe the performance of the trained model.

```
[24]: fig, axes = plt.subplots(n_rows, n_cols, figsize=(30, 32))
      axes = axes.ravel()

      for index in range(n_rows * n_cols):

          random_index = np.random.randint(0, len(X_test_gray_standardized))
          axes[index].imshow(X_test_gray_standardized[random_index].squeeze())
          axes[index].set_title(f"Ground truth: {y_test[random_index]}\nPrediction: ⬇
          ↳{test_preds[random_index]}")
```

```
axes[index].axis('off')
```



As can be seen, the trained model performs well even on the test data!

6 That is it for this project! Thank you for following up!